



SYSMON HUNTER

Detection Engineering & Analyst Console

User Manual

Version 0.3.4

Real-time Sysmon detection, correlation, and incident triage



Contents

1. Introduction

2. System Architecture

2.1 Event Pipeline

2.2 Data Flow

3. Detection Engine

3.1 Rule-Based Detection

3.2 Process-Tree Correlation

3.3 Statistical Beacon Detection

3.4 Reconnaissance-Burst Detection

3.5 Ransomware Detection

3.6 Behavior Profiling & Derived Titles

3.7 Incident Scoring

3.9 Sigma Rule Import

3.10 STIX 2.1 Export

3.11 ATT&CK Coverage Report & Navigator Export

3.12 False-Positive Similarity

4. Detection Rule Catalog

5. The Analyst Console

5.1 Incident Queue

5.2 Incident Detail Views

5.3 Explore View

5.4 Search

5.5 IOC Enrichment

5.6 ATT&CK Technique Reference

5.7 Analyst Notes

5.8 PDF Reports & STIX Export

5.9 Incident Triage

5.10 Sigma Rule Import

5.11 Theme & Database Reset

5.12 Live WebSocket Feed

5.13 ATT&CK Coverage Download

6. REST API Reference

7. Configuration Reference

8. Installation & Getting Started

9. Testing

10. Docker Deployment

11. Design Decisions

12. Project Layout



1. Introduction

Sysmon Hunter is a real-time detection and correlation engine for Windows Sysmon telemetry, paired with a live analyst console. It ingests events from an endpoint, matches them against ATT&CK-mapped detection rules, reconstructs the process tree to correlate related detections into incidents, detects command-and-control beaoning and ransomware activity statistically, and enriches indicators against external reputation sources -- all streamed to the console in real time.

The project's guiding principle is that a single matched rule is a lead, not a verdict. A lone "PowerShell ran an encoded command" event is worth a glance. The same detection sitting under a Microsoft Word process, next to a reconnaissance burst and an outbound beacon, with a file hash confirmed malicious by VirusTotal, is an incident an analyst can act on immediately. Everything documented in this manual serves that one idea: turning individually weak signals into a correlated, prioritized, investigable incident.

This manual documents the detection engine, the analyst console, the REST and WebSocket API, the configuration surface, and the deployment options as of the version printed on the cover page.

Who this is for

Detection engineers writing or tuning Sysmon-based rules; SOC analysts triaging incidents in the console; and anyone standing up the engine against their own endpoint telemetry or a corpus of recorded Sysmon events.

2. System Architecture

2.1 Event Pipeline

Every event, regardless of source -- a live Winlogbeat feed, a replayed .evtx file, or a seed script -- takes the same path through the engine. The pipeline is independent of HTTP: it runs identically under the `/ingest` endpoint, the EVTX replay script, and the automated test suite with no server running at all.

Step	Stage	What happens
1	Ingest	POST <code>/ingest</code> receives one normalized event and hands it to the pipeline. Holds no logic of its own.
2	Normalize	Winlogbeat/Sysmon JSON is converted into a uniform internal Event, so downstream code never branches on source format.
3	Observe	Every event feeds the in-memory process tree, not just the ones that match a rule -- a malicious process's ancestors are usually benign and must still be recorded.
4	Detect	The event is run against three parallel detectors: the YAML rule engine, the statistical beacon detector, and the reconnaissance-burst detector.
5	Correlate	New detections are grouped with existing ones that share a process-tree root within the correlation window, forming or extending an incident.
6	Persist & Broadcast	The detection and incident are written to SQLite and pushed to every connected console over the WebSocket feed.



2.2 Data Flow

The process tree is keyed on Sysmon's ProcessGuid, never on process ID. Windows recycles PIDs aggressively, so a tree keyed on PID would graft an unrelated later process onto a malicious parent's identity the moment that PID is reused; ProcessGuid is unique across reboots and PID reuse and does not have this failure mode.

Incidents are not simply a rule match logged to a table -- the correlator groups every detection that shares a process-tree root within a configurable time window into a single incident, and carries the full branching tree with it, not just the linear chain from root to the triggering process.

3. Detection Engine

3.1 Rule-Based Detection

127 YAML detection rules, each mapped to one or more MITRE ATT&CK technique IDs, are indexed by Sysmon Event ID so that only relevant rules are evaluated per event. Rules use Sigma-compatible matching semantics: field/operator pairs (equals, contains, startswith, endswith, re) combined with an all/any condition, and an operator's expected value may be a single value or a list, OR'd together -- for example, matching a file write against any of several known ransomware extensions in one clause.

The full catalog, with severity and ATT&CK mapping for every rule, is in Section 4.

3.2 Process-Tree Correlation

Detections rarely stand alone in a real intrusion. The correlator reconstructs each host's process ancestry from Sysmon's ProcessGuid fields, and any detections whose processes share a tree root within the correlation window (10 minutes by default) are grouped into one incident carrying the complete branching tree -- not just the chain from root to the flagged process. A foothold that spawned several children shows every branch, whether or not each child individually fired a rule.

3.3 Statistical Beacon Detection

No single network-connection event can reveal a periodic command-and-control channel; it takes a sequence of them. The beacon detector watches outbound connections (Sysmon Event ID 3) per process and scores their timing for regularity, requiring a minimum number of callbacks before rendering a verdict. It uses the median and median absolute deviation of the intervals, not the mean and standard deviation -- a live C2 session produces occasional large outliers (the operator interacts, a callback retries late), and those outliers distort a standard-deviation score badly enough to lose the channel entirely, while the median barely moves. Jitter up to and beyond Cobalt Strike's default 37% is still caught. A configurable exclusion list keeps common legitimately-periodic processes (browsers, sync clients, chat apps) from generating noise.

3.4 Reconnaissance-Burst Detection

The discovery detector counts *distinct* ATT&CK discovery techniques observed in one process tree within a short window, not raw execution counts. A script re-running systeminfo in a loop is volume without variety and stays quiet; an operator running whoami, net, nltest, and systeminfo in quick succession is variety, and is flagged as a reconnaissance burst.



3.5 Ransomware Detection

Two rules target the file-write behavior that precedes and accompanies a ransomware encryption event, in addition to the shadow-copy-deletion and recovery-disabling command-line rules that typically precede it: a ransom note dropped to disk, matched against the near-universal “how to decrypt / restore your files” naming convention used across ransomware families, and a mass file write ending in a known ransomware encryption extension (.locked, .encrypted, .crypt, .wcry, and others). Because these detections have no process command line to show as evidence -- the interesting fact is which file was written, not what process wrote it -- the console surfaces the exact matched file path everywhere the detection's detail is shown: the incident's detection list, its timeline and process-tree popups, and the full-screen tree viewer.

3.6 Behavior Profiling & Derived Titles

Rather than leaving an analyst to read a flat list of N detections, the profiling engine turns an incident's detections into a plain-language, kill-chain-ordered narrative -- for example: “gained initial access through a phishing document, executed an obfuscated PowerShell payload, harvested credentials from LSASS, beacons to C2 every ~35 seconds.” Each incident is also given a derived title summarizing its contents at a glance in the queue -- “Phishing to reconnaissance,” “Credential access with C2,” “Ransomware preparation” -- instead of a generic “Incident #4821.”

3.7 Incident Scoring

Incident severity scoring is non-linear by design: one critical LSASS access outweighs three medium-severity suspicious-path executions, and two high-severity detections together read as critical. This mirrors how an analyst actually judges risk -- a single strong signal can outweigh several weak ones, and enough weak signals together compound into a strong one. An incident is promoted to “active” once its cumulative score crosses a configurable threshold (12, by default).

3.8 Correlation Chains & Classification

Beyond the tactic-based titles in Section 3.6, three named, rule-ID-level patterns recognise a specific multi-stage story and outrank the generic titles when they match: a **ransomware activity chain** (shadow-copy deletion together with a ransom note or an encrypted-extension write), a **credential-theft campaign** (at least two distinct credential-access rules -- LSASS access, Mimikatz, DCSync, browser/KeePass credential-database staging, Kerberoasting, or AS-REP roasting -- firing on the same incident), and an **Office-to-PowerShell infection chain** (an Office application spawning a shell, followed by a PowerShell download cradle). A matched chain sets both the incident's title and a machine-readable classification field, shown as a badge in the console queue and on the incident's full page. This is deliberately an identity mechanism, not a scoring one -- every rule referenced by a chain is itself high or critical severity, so a matched chain has already cleared the actionable threshold on the ordinary scoring scale described in Section 3.7 before it is ever classified.

3.9 Sigma Rule Import

The hand-written rule catalog in Section 4 can be extended at runtime by uploading Sigma YAML files (from the public SigmaHQ corpus or written by hand) through the console's settings menu, which posts to POST /admin/rules/import-sigma. Each document is converted independently by backend/engine/sigma_import.py into this engine's own rule schema and written under rules/imported_sigma/, then the whole rule store reloads from disk -- an import is live for the very next ingested event, no restart required.



Sigma is a considerably richer language than the matcher in Section 3.1 implements, so the importer supports a deliberately bounded subset: Windows Sysmon logsource categories mapped to the matching EventID, a single selection or several combined with a plain and / or / 1 of `x*` / all of `x*` condition, one trailing and not `<filter>` exclusion, the contains / startswith / endswith / re modifiers, and automatic translation of bare glob wildcards (`*`, `?`) into the matching operator. A Sigma rule that needs nested boolean groups, an aggregation (`count()`, `near`), or a modifier this matcher cannot reproduce (`|all`, `|base64`, `|cidr`, numeric comparisons) is rejected with the specific reason rather than approximated -- the same policy the rule loader already applies to a malformed YAML file on disk: one bad rule is reported and skipped, never silently wrong and never fatal to the batch.

3.10 STIX 2.1 Export

The counterpart to Sigma import: any incident can be downloaded as a STIX 2.1 bundle from `GET /incidents/{id}/stix` (or the “Download STIX bundle” button next to the PDF report button), for another threat-intel platform to ingest without ever needing to know this engine's own schema. `backend/engine/stix_export.py` builds an identity object for this engine, one attack-pattern per distinct ATT&CK technique the incident's detections carry, one indicator per pivotable IOC already surfaced by the key-indicators consolidation in Section 5 (a C2 destination IP or domain, a file's SHA256, a persistence registry key -- deduplicated across detections, never one indicator per detection), and a report object whose `object_refs` ties all of it together, named and described from the incident's own title and behavior-profile narrative.

Object identifiers are deterministic -- a UUIDv5 seeded on stable content (a technique ID, an indicator's own pattern string, the incident ID) rather than randomly generated -- so exporting the same incident twice, or the same technique across two different incidents, produces the same STIX ID both times. That is what lets a receiving platform de-duplicate objects across imports instead of accumulating a new copy of the same technique on every download. No relationship objects are generated between indicators and techniques: the underlying data is incident-level, not per-IOC, so asserting that a specific indicator indicates a specific technique would claim precision the source data does not have.

3.11 ATT&CK Coverage Report & Navigator Export

Every other view in this manual answers what the engine has detected; this one answers what it never could. `GET /attack/coverage` counts how many rules and statistical detectors raise each ATT&CK technique across the whole corpus, and `GET /attack/coverage/navigator` (also reachable from the settings menu as “Download ATT&CK coverage”) renders that count as a MITRE ATT&CK Navigator layer -- open it at the public Navigator and every technique is colored red through green by how many rules cover it, turning months of reactive rule-writing into a single prioritized worklist for the next rule.

The report has two honesty levels, and it says which one it is giving you. `backend/data/attack_data.json` (Section 5.6) only ever contains techniques this project already references, by design -- so on its own it can rank covered techniques against each other but can never show a technique with zero rules, since that technique was filtered out before the report ever runs. `scripts/fetch_attack.py` also writes a second, full-catalog file, `backend/data/attack_index.json` -- every non-deprecated Enterprise technique, covered or not, with just enough data (ID, name, tactics) to plot on a layer. When that file is present, `backend/engine/coverage.py` produces a true gap analysis; when it is absent, the report still runs, it is just explicitly marked `partial` in both the JSON response and the layer's own description -- a missing optional file changes how much the report can show, never whether it works.



3.12 False-Positive Similarity

Every incident an analyst marks a false positive is a labeled example of noise -- `backend/engine/noise.py` is what makes that marking pay off for the next incident, without a trained model or a mountain of data to get there. A newly-opened incident is compared against the deployment's own history of confirmed false positives on four explainable signals: which detection rules fired (the most precise statement of "this is the same pattern" available, weighted highest), which ATT&CK techniques were involved, whether the same process was the root of the tree, and how much of the process chain overlaps (weighted lowest, since two unrelated incidents can share common ancestor processes like `explorer.exe` by platform convention alone).

This is deliberately not a classifier. It produces a score -- e.g. "78% similar to incident abc123, because: same detection rules, same root process" -- from the very first false positive an analyst ever marks, and the explanation is exactly the weighted arithmetic that produced it, nothing hidden and nothing to retrain. A match at or above `HUNTER_NOISE_SIMILARITY_THRESHOLD` (default 0.6, Section 7) surfaces as a dashed "probable noise" badge in the console queue and on the incident page, naming the past incident it resembles and the specific signals that matched -- a hint for an analyst to weigh, never an automatic dismissal. Only open incidents are scored; a closed or already-dismissed incident has nothing left to warn about.



4. Detection Rule Catalog

Every rule below was written and validated against real telemetry -- either a hand-built scenario exercised end-to-end through the console, or a known-malicious .evtx sample from a public detection-engineering corpus. Each has an automated true-positive case (the rule fires on the malicious sample) and a true-negative case (it stays quiet on adjacent legitimate activity) in the test suite.

Rule ID	Event	Severity	Title	ATT&CK
SYS-001	Process Create	High	Office application spawned a command interpreter	T1566.001, T1059
SYS-002	Process Create	High	PowerShell executed an encoded command	T1059.001, T1027
SYS-003	Process Create	High	LOLBin used to download a remote payload	T1105, T1218
SYS-004	Process Create	Critical	Volume shadow copies deleted	T1490
SYS-005	Process Create	High	Script engine spawned from a browser or mail client	T1204.002, T1566
SYS-006	Process Create	Medium	Executable ran from a user-writable staging path	T1204
SYS-007	Process Create	High	System binary running from the wrong path	T1036.005
SYS-008	Process Create	Medium	rundll32 executed without a DLL argument	T1218.011
SYS-009	Process Create	High	PowerShell download cradle	T1059.001, T1105
SYS-010	Process Access	Critical	Suspicious process accessed LSASS memory	T1003.001
SYS-020	Network Connection	Medium	Scripting or LOLBin process made an outbound connection	T1071
SYS-021	Network Connection	Medium	Connection to a common C2 port from a non-browser process	T1571
SYS-030	Registry Set	High	Autorun registry key modified	T1547.001
SYS-031	Registry Set	High	Windows Defender setting disabled via registry	T1685
SYS-032	Registry Set	High	Image File Execution Options debugger set	T1546.012
SYS-034	Process Create	High	Control panel item (.cpl) executed from a user-writable path	T1218.002
SYS-035	Process Create	High	Script host executing an encoded script	T1059.005, T1059.007, T1027
SYS-036	Process Create	High	Script host spawned by rundll32	T1059, T1218.011
SYS-037	Process Create	High	rundll32 invoking a known LOLBAS DLL export	T1218.011, T1204.001
SYS-038	Process Create	High	IIS credentials extracted via appcmd	T1552.001, T1003
SYS-039	Process Create	Medium	appcmd executed by a web server worker or script host	T1552.001
SYS-040	Image Load	High	PowerShell engine loaded by a non-PowerShell process	T1059.001
SYS-041	Process Access	Critical	LSASS opened with credential-dumping access rights	T1003.001
SYS-050	File Create	High	File dropped into a Startup folder	T1547.001
SYS-051	File Create	High	Office application wrote a script or executable	T1566.001
SYS-060	Pipe Created	High	Named pipe matching a known C2 default	T1071, T1055
SYS-070	Process Create	Medium	WMI persistence staged via mofcomp or wmic	T1546.003
SYS-071	WMI Event	High	WMI event consumer registered for code execution	T1546.003



Rule ID	Event	Severity	Title	ATT&CK
SYS-072	Pipe Created	High	Named pipe matching PsExec-style remote execution	T1021.002, T1569.002
SYS-073	Process Create	High	PsExec service binary executed on this host	T1569.002, T1021.002
SYS-074	Process Create	High	Regsvr32 registering a scriptlet via scrobj.dll (Squiblydoo)	T1218.010
SYS-075	Process Create	Medium	Certutil used to decode or encode a local file	T1140
SYS-076	Process Create	Medium	BITS job created via the Start-BitsTransfer cmdlet	T1197
SYS-077	Driver Load	High	Unsigned kernel driver loaded	T1211, T1562.001
SYS-078	Registry Set	High	Registry Run key written by the WMI provider host	T1047, T1547.001
SYS-079	Process Create	High	FTP client spawned a child process	T1105, T1202
SYS-080	File Create	Critical	Ransom note dropped	T1486, T1491.001
SYS-081	File Create	Critical	File written with a known ransomware encryption extension	T1486
SYS-082	Process Create	Critical	Credential dumping via comsvcs.dll MiniDump export	T1003.001
SYS-083	Registry Set	High	Fileless UAC bypass via auto-elevate registry hijack	T1548.002
SYS-084	Process Create	High	cmstp executed with a silent or auto-install flag	T1218.003
SYS-085	Registry Set	High	Port forwarding rule added via netsh	T1090.001
SYS-086	Registry Set	High	PowerShell script block logging disabled via registry	T1562.001
SYS-087	Registry Create/Delete	High	PowerShell Constrained Language Mode lockdown policy removed	T1562.001
SYS-088	Process Create	Critical	IIS worker process spawned a command shell	T1505.003
SYS-089	Process Create	Critical	SQL Server process spawned a command shell	T1059.003, T1190
SYS-090	Process Create	Medium	Process spawned by the PowerShell remoting host	T1021.006
SYS-091	Registry Set	Medium	Local account or administrators-group membership changed in the SAM hive	T1136.001, T1098
SYS-092	Process Create	High	PE metadata claims a trusted binary, but it runs from a staging path	T1036.005
SYS-093	Process Create	High	Scheduled task created via schtasks	T1053.005
SYS-094	Registry Set	High	New service image path points at a user-writable staging location	T1543.003
SYS-095	Remote Thread Created	Critical	Remote thread created inside LSASS	T1055, T1003.001
SYS-096	Process Create	High	Firewall rule added to allow inbound traffic	T1562.004
SYS-097	Process Create	High	Windows event log cleared via wevtutil	T1070.001
SYS-098	Process Create	Medium	Archive utility invoked with a password, staging data for exfiltration	T1560.001
SYS-099	Process Create	High	.NET installer utility executed from a user-writable staging path	T1218.004, T1218.009
SYS-100	Process Create	High	mshta executed a remote HTA	T1218.005
SYS-101	Process Create	Medium	Compiled HTML Help opened from a user-writable staging path	T1218.001



Rule ID	Event	Severity	Title	ATT&CK
SYS-102	Process Create	High	msiexec installed a package from a remote URL	T1218.007
SYS-103	Process Create	Medium	Account added to the local Administrators group via net.exe	T1098, T1136.001
SYS-104	Process Create	High	A security or update service was stopped from the command line	T1489, T1562.001
SYS-105	Process Create	High	Sysinternals sdelete executed	T1070.004
SYS-106	Process Create	Medium	cipher used to wipe free disk space	T1070.004
SYS-107	Registry Set	Medium	Remote Desktop enabled via registry	T1021.001
SYS-108	Process Create	Critical	Credential hive saved to disk via reg.exe	T1003.002
SYS-109	Process Create	Critical	Sysinternals procdump run against LSASS	T1003.001
SYS-110	Process Create	Critical	NTDS.dit extracted via ntdsutil	T1003.003
SYS-111	Process Create	High	mavinject used to inject into a running process	T1055.001
SYS-112	Process Create	Medium	WSL used to execute a command outside its Linux filesystem	T1202, T1059.004
SYS-113	Process Create	High	MSBuild ran a project file staged in a user-writable path	T1127.001
SYS-114	Process Create	Medium	Uncommon signed proxy-execution binary run against a staged target	T1218, T1216
SYS-115	Process Create	High	RDP session hijacked via tscon	T1563.002
SYS-116	Process Create	High	Active Directory reconnaissance tool executed	T1087, T1482
SYS-117	Process Create	Medium	Cloud sync/transfer tool executed	T1567.002
SYS-118	Process Create	High	Kerberoasting tooling executed	T1558.003
SYS-119	Process Create	High	PowerShell launched in version-2 downgrade mode	T1059.001
SYS-120	Process Create	High	Command line references a known AMSI-bypass signature	T1562.001
SYS-121	Process Create	Medium	bitsadmin used to transfer a file	T1197
SYS-122	Registry Set	High	CLSID InprocServer32 handler registered under a user's own hive	T1546.015
SYS-123	Process Create	High	Command line references a known privilege-escalation token-theft tool	T1134
SYS-124	22	Medium	DNS query to a dynamic DNS domain	T1071.004, T1568
SYS-125	23	Medium	Executable deleted from a staging path shortly after being dropped	T1070.004
SYS-126	25	Critical	Process image tampering detected (hollowing or doppelganging)	T1055.012, T1055.013
SYS-127	15	High	Executable content staged inside an NTFS alternate data stream	T1564.004
SYS-128	9	High	Raw volume access by a process outside known disk utilities	T1006
SYS-129	2	Medium	File creation timestamp changed outside a servicing process	T1070.006
SYS-130	Process Create	Critical	DCSync-style directory replication requested from a command line	T1003.006



Rule ID	Event	Severity	Title	ATT&CK
SYS-131	Process Create	Critical	Command line references a known Mimikatz module or command	T1003.001
SYS-132	Process Create	High	Odbcconf used as a signed proxy-execution binary	T1218.008
SYS-133	Process Create	High	mmc.exe loaded a Management Saved Console from a staging path	T1218.014
SYS-135	File Create	High	Non-browser process touched a browser credential database	T1555.003
SYS-136	File Create	High	Script interpreter touched a KeePass database	T1555
SYS-137	File Create	High	Browser credential database written into a staging directory	T1555.003
SYS-138	Process Create	Medium	Network configuration discovery command executed	T1016
SYS-139	Process Create	Medium	Active network connections enumerated via netstat	T1049
SYS-140	Process Create	Medium	Recursive filesystem enumeration executed	T1083
SYS-141	Process Create	High	Security or EDR product enumerated from the command line	T1518.001
SYS-142	Process Create	Medium	Archive created with its header/file-list encrypted	T1560.001
SYS-143	Process Create	High	Rclone ran a copy/sync against a remote destination	T1567.002
SYS-144	Process Create	High	Command-line tooling uploaded a file to a remote server	T1048
SYS-148	Process Create	High	SharpHound invoked via script or named by its collection arguments	T1482
SYS-149	Process Create	Medium	Active Directory account or group enumeration command executed	T1087
SYS-150	Process Create	High	AS-REP roasting tooling executed	T1558.004
SYS-151	Process Create	High	ClickFix/FileFix decoy verification lure pasted from Explorer	T1204.004
SYS-152	Process Create	High	PowerShell launched from Explorer with bypass + hidden window	T1204.004
SYS-153	Process Create	High	Script host from Explorer immediately fetching remote code	T1204.004
SYS-154	Process Create	High	PowerShell reads and executes clipboard contents	T1204.004
SYS-155	Process Create	High	Remote-access/RMM software launched from a non-Explorer parent	T1219
SYS-156	Process Create	High	Remote-access/RMM software installed with a silent flag	T1219
SYS-157	Process Create	High	PowerShell reads and overwrites the clipboard (crypto-clipper)	T1115
SYS-158	Process Create	High	Problem Steps Recorder invoked for silent screen capture	T1113
SYS-159	Process Create	High	Chat-service webhook or bot API used as a covert channel	T1102
SYS-160	Process Create	High	Public paste service used as a C2 dead-drop	T1102
SYS-161	Process Create	Medium	Archive utility packaged a broad user data directory	T1074



Rule ID	Event	Severity	Title	ATT&CK
SYS-162	Process Create	Medium	robocopy mass-mirrored a broad user directory tree	T1119, T1074
SYS-163	Process Create	Critical	Rubeus used to pass or renew a Kerberos ticket	T1550.003
SYS-164	Process Create	Critical	Mimikatz used to pass a hash or a Kerberos ticket	T1550.002, T1550.003
SYS-165	Process Create	Critical	Known Group Policy abuse tooling or cmdlet invoked	T1484.001
SYS-166	Process Create	Critical	Domain trust modified to disable SID-filtering protections	T1484.002
SYS-167	Process Create	Critical	Command-line destructive wipe of a broad user directory or volume	T1485
SYS-168	Process Create	Critical	PowerShell bulk-disabled or reset a broad set of AD accounts	T1531
SYS-169	File Create	Critical	Script interpreter staged a cloud CLI credential file	T1552.001
SYS-170	Process Create	Critical	Cloud instance metadata service queried by a non-agent process	T1552.005
SYS-171	File Create	High	Script interpreter staged an SSH or PuTTY private key	T1552.004
SYS-172	Process Create	High	WMI used to create a process via Win32_Process	T1047
SYS-173	Process Create	High	mmc.exe spawned a shell consistent with MMC20.Application DCOM abuse	T1021.003
SYS-174	Process Create	High	Docker container launched with --privileged	T1611

Rule IDs follow no severity ordering; they are assigned sequentially as rules are added. "Event" is the Sysmon Event ID the rule is indexed under (Process Create = 1, Network Connection = 3, Driver Load = 6, Image Load = 7, Remote Thread Created = 8, Process Access = 10, File Create = 11, Registry Create/Delete = 12, Registry Set = 13, Pipe Created = 17, WMI Event = 20).



5. The Analyst Console

The console is a single-page, dark-surface application served directly by the backend at the site root, with a live WebSocket feed so every connected analyst sees the same state update in real time.

5.1 Incident Queue

The primary view: every incident the engine has correlated, most recent first, filterable between “Triage” (actionable incidents only, i.e. those over the score threshold), “All” (everything the engine is tracking, including sub-threshold activity it has not promoted), and “Closed” (incidents an analyst has already resolved -- see Section 5.9). Each row shows the derived title, severity, host, cumulative score, the process chain, and ATT&CK technique chips at a glance, before any drill-down. An incident that matches one of the correlation chains from Section 3.8 (ransomware, credential-theft campaign, Office-to-PowerShell) additionally carries a filled classification badge next to its title, repeated on the incident's full page.

5.2 Incident Detail Views

Expanding an incident reveals its behavior-profile narrative, then three interchangeable views of the same underlying data:

- **List** -- every detection with its full forensic context: process, parent, user, integrity level, hashes (with a one-click VirusTotal lookup), and, for file-write detections, the exact matched file path.
- **Timeline** -- detections placed in sequence with the real elapsed time between steps; clicking a node opens a floating forensic detail popup.
- **Process tree** -- the complete branching tree the incident spans, not just the chain to the triggering process. Nodes that fired a detection are colour-coded by severity; benign context processes are shown hollow so the whole tree's shape is visible, not only the flagged parts.

5.3 Explore View

The inline detail views are deliberately compact so they fit beside the detection list, which is too little room for a wide tree, a long timeline, or a long log list. “Explore” hands the same incident to a dedicated full-screen page (`/incident/{id}/explore`) with three tabs switchable in place, each a click away from the other: process tree, timeline, and a plain scrollable log. The tree and timeline gain click-and-drag panning, scroll-wheel zoom centred on the cursor, a zoom-to-fit control, and keyboard shortcuts (+ / - / 0). Clicking a node or a timeline entry opens the same forensic detail panel used inline.

5.4 Search

One search box combines free text with field filters, mixed freely in a single query -- for example, `powershell severity:critical host:fin-ws`. Free text matches anywhere an analyst might remember something about an incident: a process name, a command-line fragment, a hash, a rule ID, the title. Field filters narrow precisely by substring match, not exact match, so `host:fin-ws` finds `FIN-WS-07`.

Filter	Matches
<code>host:</code>	Substring match against the incident's host.
<code>severity:</code>	Exact match: info, low, medium, high, or critical.
<code>technique:</code>	Matches an ATT&CK technique ID, exact or as a prefix.



Filter	Matches
rule:	Matches a detection rule ID.
user:	Matches the forensic user field on a detection.
command_line:	Scoped substring match against a detection's command line only.
actionable:	true/false -- restricts to incidents over (or under) the score threshold.

5.5 IOC Enrichment

IP addresses, domains, and file hashes can be checked on demand against AbuseIPDB (IP reputation) and VirusTotal (IP, domain, and file-hash reputation). Results are cached for an hour. Every provider degrades gracefully to “unavailable” with no API key configured -- enrichment never blocks the rest of the console from working. Private IP ranges are never sent to a third party, and the strongest available hash (SHA-256 over SHA-1 over MD5) is always preferred, since MD5 collisions are cheap enough for malware authors to produce deliberately.

5.6 ATT&CK Technique Reference

Every technique chip shown anywhere in the console is clickable and opens its official MITRE description -- name, tactics, and full text -- from a local copy of the ATT&CK STIX dataset, with a link out to the ATT&CK site. No lookup leaves the console for an analyst who does not have T1003.001 memorized.

5.7 Analyst Notes

Each incident's full-page view carries a free-text analyst note, capped at 500 words -- a summary, not a report -- with autosave on the incident page and a live word counter.

5.8 PDF Reports & STIX Export

Every incident can be downloaded as a self-contained PDF report, suitable for attaching to a ticket or handing to an incident-response lead. It is not a screenshot of the console -- it is the investigation written down: the behavior-profile narrative, the process chain, every detection with its forensics, and the incident's key indicators. Generated server-side with reportlab, requiring no headless browser in the deployment image. The button beside it, “Download STIX bundle”, exports the same incident as a machine-readable STIX 2.1 bundle instead -- see Section 3.10 for exactly what it contains.

5.9 Incident Triage

Every incident carries a status: open, closed, or a verdict such as false positive or benign, set from a “Set verdict” menu on the incident card and cleared with a “Close” action. Any state can transition to any other -- a closed incident can be reopened if new evidence surfaces -- and there is no workflow enforced beyond that; the analyst is the one doing the triage, not the console. Closed incidents drop out of the default queue view and are reachable from the “Closed” filter tab (Section 5.1), so a shift's queue only ever shows what still needs a decision.

Once at least one incident has been marked a false positive, later open incidents that resemble it closely enough gain a dashed “probable noise” badge next to the status badge, naming the similarity percentage -- see Section 3.12 for the scoring behind it.



5.10 Sigma Rule Import

The same settings menu that holds the theme toggle also holds “Import Sigma rules”, which opens a file picker accepting one or more .yaml/.yml files. Selecting files immediately uploads them to `POST /admin/rules/import-sigma`; a report modal then lists every document that was accepted (with its generated rule ID) and every one that was rejected (with the specific reason), so a partially-successful batch is never a silent partial success. See Section 3.9 for exactly what subset of Sigma is supported.

5.11 Theme & Database Reset

A settings menu in the console header holds a light/dark theme toggle and one destructive action: wiping the database. The reset is confirmed with a dialog, since there is no undo -- every detection and incident is permanently removed, on every connected console, via the same WebSocket broadcast that delivers live updates.

5.12 Live WebSocket Feed

The console holds a persistent WebSocket connection and reconnects automatically on drop, so it can be left open on a wall display indefinitely. New detections, incident updates, and database resets are pushed to every connected client the instant they happen -- there is no polling and no manual refresh.

5.13 ATT&CK Coverage Download

The same settings menu holds “Download ATT&CK coverage”, a direct link to `GET /attack/coverage/navigator` -- no upload step, no report modal, just a MITRE ATT&CK Navigator layer file ready to open at the public Navigator. See Section 3.11 for what it contains and how it degrades when the full technique index is not present.



6. REST API Reference

The full interactive OpenAPI documentation is served at `/api/docs` (Swagger UI) and `/api/redoc` (ReDoc) on a running instance. The table below is a static summary of every route.

Method	Path	Description
POST	<code>/ingest</code>	Accept one normalized Sysmon/Winlogbeat event; runs it through the full pipeline.
GET	<code>/incidents</code>	List incidents, most recent first.
GET	<code>/incidents/{id}</code>	Full incident detail: detections, process tree, forensics.
PUT	<code>/incidents/{id}/status</code>	Close an incident, mark it a false positive, or reopen it.
PUT	<code>/incidents/{id}/notes</code>	Set the incident's analyst note (500-word limit).
DELETE	<code>/incidents/{id}/notes</code>	Clear the incident's analyst note.
GET	<code>/incidents/{id}/profile</code>	Kill-chain narrative summary for the incident.
GET	<code>/incidents/{id}/report</code>	Generate and download the incident's PDF report.
GET	<code>/incidents/{id}/stix</code>	Generate and download the incident as a STIX 2.1 bundle.
GET	<code>/detections</code>	List raw detections, most recent first.
GET	<code>/search</code>	Free-text and field-filtered search across incidents.
GET	<code>/attack/coverage</code>	Rule-coverage report: rule/detector count per ATT&CK technique.
GET	<code>/attack/coverage/navigator</code>	The coverage report as a downloadable MITRE ATT&CK Navigator layer.
GET	<code>/attack/{technique_id}</code>	MITRE ATT&CK technique description, from the local dataset.
GET	<code>/enrich</code>	On-demand reputation lookup for an IP, domain, or hash.
DELETE	<code>/admin/database</code>	Wipe all detections and incidents; reset the live engine.
POST	<code>/admin/rules/import-sigma</code>	Convert and load one or more Sigma YAML files; live immediately.
GET	<code>/health</code>	Liveness/readiness probe: rule count, tracked processes.
GET	<code>/metrics</code>	Prometheus text-exposition scrape target: request, ingest, and detection counters.
WS	<code>/ws</code>	Live event stream: new detections, incident updates, resets.
GET	<code>/</code>	The analyst console (single page app).
GET	<code>/incident/{id}</code>	Full-page view of one incident, with notes editor.
GET	<code>/incident/{id}/explore</code>	Full-screen Explore view: process tree, timeline, or logs, picked by tab.
GET	<code>/incident/{id}/tree</code>	Alias for <code>/explore?view=tree</code> , kept for old links.



7. Configuration Reference

Every tunable lives in one settings object and can be overridden from the environment or a `.env` file, without touching code. All variables are prefixed `HUNTER_`.

Variable	Default	Purpose
HUNTER_DB_URL	sqlite+aiosqlite:///data/hunter.db	Database connection string.
HUNTER_API_KEY	(unset)	Shared secret for the JSON API. Unset = no auth (trusted network).
HUNTER_CORRELATION_WINDOW_MINUTES	10	Time window for grouping detections into one incident.
HUNTER_INCIDENT_SCORE_THRESHOLD	12	Cumulative severity score to promote an incident to active.
HUNTER_PROCESS_TTL_MINUTES	120	How long an inactive process stays in the in-memory tree.
HUNTER_BEACON_ENABLED	true	Enable/disable statistical beacon detection.
HUNTER_BEACON_MIN_CONNECTIONS	6	Callbacks required before a beacon verdict.
HUNTER_BEACON_REGULARITY_THRESHOLD	0.75	Minimum regularity score (0-1) to flag a beacon.
HUNTER_BEACON_MIN_INTERVAL_SECONDS	5.0	Fastest interval considered beacon-like, not streaming.
HUNTER_BEACON_MAX_INTERVAL_SECONDS	3600.0	Slowest interval still distinguishable from a scheduled task.
HUNTER_DISCOVERY_ENABLED	true	Enable/disable reconnaissance-burst detection.
HUNTER_DISCOVERY_MIN_DISTINCT	4	Distinct recon techniques required to flag a burst.
HUNTER_NOISE_SIMILARITY_THRESHOLD	0.6	Similarity score (0-1) an open incident needs to be flagged as probable noise.
HUNTER_ABUSEIPDB_API_KEY	(unset)	Optional key for IP reputation enrichment.
HUNTER_VIRUSTOTAL_API_KEY	(unset)	Optional key for hash/IP/domain reputation enrichment.
HUNTER_ENRICHMENT_CACHE_TTL_SECONDS	3600	How long a reputation lookup is cached.
HUNTER_LOG_JSON	false	One JSON object per log line, for log aggregators, instead of human-readable text.
HUNTER_INGEST_RATE_LIMIT_PER_SECOND	0	Token-bucket rate limit on <code>/ingest</code> per source IP. 0 disables it.
HUNTER_INGEST_RATE_LIMIT_BURST	50	Burst allowance above the steady-state <code>/ingest</code> rate limit.



8. Installation & Getting Started

Requires Python 3.11 or newer.

```
python -m venv .venv
source .venv/bin/activate    # Windows: .venv\Scripts\Activate.ps1
pip install -r requirements.txt

python -m alembic upgrade head    # create the database schema
python -m uvicorn backend.main:app --reload --host 0.0.0.0 --port 8000
```

Console: <http://localhost:8000> · API docs: <http://localhost:8000/api/docs>

Seeding demo data

```
python scripts/seed_apr.py    # one deep multi-stage intrusion
python scripts/seed_demo.py   # varied incidents across the kill chain
python scripts/seed_rw.py     # a full ransomware chain in one incident
```

Analysing a real sample

Replay a Sysmon .evtx file -- from a lab VM, or a public corpus such as EVTX-ATTACK-SAMPLES, where each file exercises one ATT&CK technique:

```
pip install evtx
python scripts/replay_evtx.py --file samples/sysmon_credential_access.evtx
```

Optional: IOC enrichment API keys

Works with no keys configured (every provider reports unavailable). For live reputation data, add free-tier API keys to a .env file:

```
HUNTER_ABUSEIPDB_API_KEY=...
HUNTER_VIRUSTOTAL_API_KEY=...
```

Optional: API key (authentication)

The server binds to all interfaces (0.0.0.0) by default, so it is reachable from anywhere on the network it runs on -- fine for a single trusted network, not fine once it is reachable more broadly. Setting HUNTER_API_KEY makes every JSON endpoint (not the console's own HTML pages) require a matching X-API-Key header; see backend/api/auth.py.

Generate a random secret with whichever tool is on hand:

```
python -c "import secrets; print(secrets.token_hex(32))" # any platform
openssl rand -hex 32                                     # Linux / macOS
```

Then set it in .env:

```
HUNTER_API_KEY=<paste the generated string here>
```

The console detects a 401 on its own and prompts once for the key, then remembers it in the browser -- no other setup needed. Treat the key like a password: do not commit it, and prefer HTTPS (e.g. behind a reverse proxy) if the server is reachable outside a trusted network.

9. Testing

```
pip install pytest pytest-asyncio
python -m pytest    # 598 tests
```



The test suite doubles as documentation: each design decision named in Section 11 has a test named after it, and every shipped detection rule is validated against a true positive it must catch and a true negative it must ignore.

10. Docker Deployment

```
docker compose up --build
```

A multi-stage build keeps the runtime image to the application and its virtual environment only -- the build tooling is discarded. The container runs as an unprivileged user, applies database migrations automatically on start, and exposes a health check at `/health`.

11. Design Decisions

The choices below are what separate this engine from a pattern match over a log file. Each is enforced by an automated test named after it.

Correlation keys on ProcessGuid, never PID.

Windows recycles process IDs aggressively; a PID-keyed tree would graft a malicious child onto whatever unrelated process later inherits its parent's PID. ProcessGuid is unique across reboots and PID reuse.

The process tree observes everything, not just detections.

A malicious process's ancestors are usually entirely benign, so the process that never triggers a rule on its own must still be recorded, or it can never appear as the root of a phishing chain.

Incident scoring is non-linear.

One critical LSASS access outweighs three medium-severity suspicious-path executions. Two high-severity detections read as critical together. An incident can be worse than any single rule that composes it.

Beaconing uses median and median absolute deviation, not mean and standard deviation.

A live C2 session produces outliers constantly. One unusually long gap in an otherwise 60-second beacon wrecks a standard-deviation-based score and loses the channel; the median barely notices. Jitter up to Cobalt Strike's default 37% is still caught.

Discovery counts distinct techniques, not raw execution counts.

A script re-running one command in a loop is volume without variety. An operator running several different reconnaissance commands in quick succession is variety, and only the second pattern is flagged.

Enrichment works with no API keys, and never leaks internal data.

Every provider degrades to "unavailable" without a key rather than failing the request. Private IP ranges are never sent to a third party. The strongest available hash is always preferred over a weaker one.

A missing optional data file degrades the ATT&CK coverage report, never crashes it.

Without the full technique index the coverage report and its Navigator export still run -- they just say so, falling back to ranking only the techniques this project already covers instead of refusing to answer at all.



12. Project Layout

```
sysmon-hunter/  
+-- backend/  
| +-- main.py          FastAPI app, lifespan, background sweep, /metrics  
| +-- config.py        all tunables  
| +-- logging_setup.py text/JSON log formatting  
| +-- api/             ingest, detections, incidents, attack, enrich,  
| |                   report, search, notes, admin, ws, serializers,  
| |                   rate_limit  
| +-- engine/          normalizer, rule_loader, matcher, correlator,  
| |                   beacon, discovery, attack, coverage, enrichment,  
| |                   search, report, profile, pipeline, sigma_import,  
| |                   stix_export, noise, metrics  
| +-- models/          schemas, db  
| `-- data/            attack_data.json (ATT&CK technique lookup),  
|                   attack_index.json (full catalog, coverage gaps)  
+-- rules/             127 YAML detection rules, by Event ID, plus  
|                   imported_sigma/ for runtime Sigma imports  
+-- frontend/          console.html, incident.html, tree.html, static/{css,js}  
+-- migrations/        Alembic  
+-- scripts/           seed_apr, seed_demo, seed_rw, seed_full_coverage,  
|                   replay_evtx, fetch_attack  
+-- docs/              screenshots  
`-- tests/             598 tests
```

Built as a hands-on threat-detection research project. The engine targets a single collector; scaling to many would mean moving the queue to Redis and the store to Postgres, both isolated behind small interfaces so that change touches one file each.